

Course of Web Information Retrieval

Homework 3

Simone Agostinelli 1523559, Giacomo Vettraino 1594722

1 Ham/Spam Classifier

1.1 KNN Classifier

1.1.1 Dataset

Our dataset is composed by 347 text files representing ham and spam comments related to YouTube videos. These files are composed according to the following table:

	Training Set	Test Set
Spam	122	53
Ham	120	52

1.1.2 Vectorizer and Classifier Parameters

Firstly we consider the process of vectorization that we have to apply in order to transform our text files into a matrix of TfIdf features. For this step we have:

tokenizer $\in \{None, stemming_tokenizer, stemming_tokenizer_stopwords_filter\}$
override the string tokenization function if its value is different from *None*;

ngram_range $\in \{(1, 1), (1, 2), (1, 3)\}$
specify the lower and upper boundary of the range of n-values for different n-grams to be extracted;

While, for the classification phase we have:

n_neighbors $\in \{1, 3, 5, 7, 9\}$

specify the number of neighbors to use for classifying points;

weights $\in \{ "uniform", "distance" \}$

specify the weight function in prediction;

1.1.3 Training-Validation Phase

Training and validation are performed by using the sklearn Python function **GridSearchCV** in an automated fashion. Indeed, it performs an exhaustive search of the best parameters values configuration by fitting the specific model (or a series of models transformation if pipeline is used); furthermore, it can output the best estimator and the best parameters configuration by computing scores on this combinations using a scoring function chosen by the user.

In particular, in our case, firstly we decided to use **Pipeline** sklearn function, in order to construct the different steps that we want to cross-validate: vectorization and classifier fitting. **Pipeline** function has input parameter:

steps

List of (name, transform) tuples that are chained, in the order in which they are chained, with the last object an estimator.

```
1 | vectorizer = TfidfVectorizer(strip_accents= None ,
2 | preprocessor = None ,)
3 | knn = KNeighborsClassifier()
4 | pipeline = Pipeline([('vect', vectorizer), ('knn', knn),])
```

Then, we can directly make use of **GridSearchCV** function, in order to setup the automated search for the best parameter configuration.

GridSearchCV has as input parameters:

estimator

This is the estimator on which we want to perform our grid search; in this case we set it to the pipeline just created, as its last step represent an estimator, our KNN Classifier;

param_grid

Dictionary with parameters names (string) as keys and lists of parameter settings to try as values, or a list of such dictionaries, in which case the grids spanned by each dictionary in the list are explored; this enables searching over any sequence of parameter settings;

scoring

A string (see model evaluation documentation) or a scorer callable object / function with signature `scorer(estimator, X, y)`; in this case Matthews correlation coefficient is used:

$$|MCC| = \frac{TP \times TN - FP \times FN}{\sqrt{(TP + FP)(TP + FN)(TN + FP)(TN + FN)}}$$

n_jobs

Number of jobs to run in parallel, in this way we can control the parallelism of our program, in order to speed up it running time;

cv

determines the cross-validation splitting strategy; an integer specify the number of folds in a KFold, 10 in our case;

```

1 | parameters = {
2 | 'vect__tokenizer': [None, stemming_tokenizer,
3 | stemming_tokenizer_stopwords_filter],
4 | 'vect__ngram_range': [(1, 1), (1, 2), (1, 3)],
5 | 'knn__n_neighbors': [1, 3, 5, 7, 9],
6 | 'knn__weights': ["uniform", "distance"]
7 | }
8 | grid_search = GridSearchCV(
9 | pipeline,
10 | parameters,
11 | scoring = metrics.make_scorer(metrics.matthews_corrcoef),
12 | cv = 10,
13 | n_jobs = 8)

```

1.1.4 Best Parameters Values

After performing cross-validation, as explained before, we have found that the best parameters configuration is:

Parameter	Value
knn__n_neighbors	9
knn__n_weights	distance
vect__ngram_range	(1,2)
vect__tokenizer	None

1.1.5 Results for Classification

Output of `metrics.classification_report`:

1	-----				
2		precision	recall	f1-score	support
3					
4	Ham	0.83	0.96	0.89	52
5	Spam	0.96	0.81	0.88	53
6					
7	avg / total	0.90	0.89	0.89	105
8	-----				

where:

$$\text{precision} = \frac{TP}{TP+FP}$$

$$\text{recall} = \frac{TP}{TP+FN}$$

$$\text{f1 score} = \frac{2(\text{precision} \times \text{recall})}{(\text{precision} + \text{recall})}$$

support is the number of samples belonging to that class

The Confusion Matrix:

	Predicted-Ham	Predicted-Spam
True-Ham	50	2
True-Spam	10	43

The **Normalized-Accuracy** value: 0.885714285714

The **Matthews Correlation Coefficient**: 0.780832919631

2 Sentiment Analysis

2.1 KNN Classifier

2.1.1 Dataset

Our dataset is composed by 1115 text files representing positive and negative sentences. These files are composed according to the following table:

	Training Set	Test Set
Negative	249	250
Positive	308	308

2.1.2 Vectorizer and Classifier Parameters

Firstly we consider the process of vectorization that we have to apply in order to transform our text files into a matrix of Tfidf features. For this step we have:

tokenizer $\in \{None, stemming_tokenizer, stemming_tokenizer_stopwords_filter\}$
 override the string tokenization function if its value is different from *None*;

ngram_range $\in \{(1, 1), (1, 2), (1, 3)\}$
 specify the lower and upper boundary of the range of n-values for different n-grams to be extracted;

While, for the classification phase we have:

n_neighbors $\in \{1, 3, 5, 7, 9\}$

specify the number of neighbors to use for classifying points;

weights $\in \{uniform, distance\}$

specify the weight function in prediction ;

2.1.3 Training-Validation Phase

See subsection 1.1.3

2.1.4 Best Parameters Values

After performing cross-validation, as explained before, we have found that the best parameters configuration is:

Parameter	Value
knn__n_neighbors	3
knn__n_weights	distance
vect__ngram_range	(1,2)
vect__tokenizer	stemming_tokenizer_stopwords_filter

2.1.5 Results for Classification

Output of **metrics.classification_report**:

1	-----
2	
3	
4	Positive 0.85 0.94 0.89 308
5	negative 0.92 0.80 0.85 250
6	
7	avg / total 0.88 0.88 0.88 558
8	-----

The **Confusion Matrix**:

	Predicted-Positive	Predicted-Negative
True-Positive	290	18
True-Negative	51	199

The **Normalized-Accuracy** value: 0.876344086022

The **Matthews Correlation Coefficient**: 0.752375671874

2.2 MultinomialNB Classifier

2.2.1 Dataset

Our dataset is composed by 1115 text files representing positive and negative sentences. These files are composed according to the following table:

	Training Set	Test Set
Negative	249	250
Positive	308	308

2.2.2 Vectorizer and Classifier Parameters

Firstly we consider the process of vectorization that we have to apply in order to transform our text files into a matrix of Tfidf features. For this step we have:

tokenizer $\in \{None, stemming_tokenizer, stemming_tokenizer_stopwords_filter\}$
 override the string tokenization function if its value is different from *None*;

ngram_range $\in \{(1, 1), (1, 2), (1, 3)\}$
 specify the lower and upper boundary of the range of n-values for different n-grams to be extracted;

While, for the classification phase we have:

alpha $\in \{0.001, 0.01, 1, 10\}$ specify the additive smoothing parameter;

2.2.3 Training-Validation Phase

See subsection 1.1.3

2.2.4 Best Parameters Values

After performing cross-validation, as explained before, we have found that the best parameters configuration is:

Parameter	Value
mnbc__alpha	1
vect__ngram_range	(1,1)
vect__tokenizer	stemming_tokenizer_stopwords_filter

2.2.5 Results for Classification

Output of `metrics.classification_report`:

1	-----
2	
3	
4	precision recall f1-score support
5	Positive 0.90 0.98 0.94 308
6	negative 0.98 0.86 0.92 250
7	avg / total 0.93 0.93 0.93 558
8	-----

The **Confusion Matrix**:

	Predicted-Positive	Predicted-Negative
True-Positive	303	5
True-Negative	34	216

The **Normalized-Accuracy** value: 0.930107526882
The **Matthews Correlation Coefficient**: 0.862006201146

2.3 Linear SVC

2.3.1 Dataset

Our dataset is composed by 1115 text files representing positive and negative sentences. These files are composed according to the following table:

	Training Set	Test Set
Negative	249	250
Positive	308	308

2.3.2 Vectorizer and Classifier Parameters

Firstly we consider the process of vectorization that we have to apply in order to transform our text files into a matrix of Tfidf features. For this step we have:

tokenizer $\in \{None, stemming_tokenizer, stemming_tokenizer_stopwords_filter\}$
override the string tokenization function if its value is different from *None*;

ngram_range $\in \{(1, 1), (1, 2), (1, 3)\}$
specify the lower and upper boundary of the range of n-values for different n-grams to be extracted;

While, for the classification phase we have:

C $\in \{0.01, 0.1, 1.0, 10.0, 100.0\}$
specify the penalty parameter C value of the error term ;

2.3.3 Training-Validation Phase

See subsection 1.1.3

2.3.4 Best Parameters Values

After performing cross-validation, as explained before, we have found that the best parameters configuration is:

Parameter	Value
svc__C	10
vect__ngram_range	(1,3)
vect__tokenizer	stemming_tokenizer_stopwords_filter

2.3.5 Results for Classification

Output of `metrics.classification_report`:

1	-----				
2		precision	recall	f1-score	support
3					
4	Positive	0.96	0.97	0.97	308
5	negative	0.97	0.95	0.96	250
6					
7	avg / total	0.96	0.96	0.96	558
8	-----				

The Confusion Matrix:

	Predicted-Positive	Predicted-Negative
True-Positive	300	8
True-Negative	13	237

The **Normalized-Accuracy** value: 0.962365591398

The Matthews Correlation Coefficient: 0.923917742518